

AYMANE AIT DADS

AI Intern 2026 | Graph AI · LLM Fine-tuning · ML Pipelines

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 11, 2026

Hiring Team - AI Intern 2026

Arango · Sophia Antipolis, France

Dear Arango Hiring Team,

Arango's core thesis - that LLMs fail without a **trusted data foundation** grounded in **contextual structure** - is exactly the problem I ran into when building production AI systems. Fine-tuning **CLIP ViT-L/14 (428M parameters)** on 250K samples spanning 25+ generative model types taught me that raw model power means nothing without the right representational context: the difference between #1 on the EURECOM leaderboard and mid-table was entirely about what structure you fed the model, not the model itself.

At Orange Maroc, I engineered an **end-to-end ML pipeline** over **100K+ telemetry records** - handling raw ingestion, KMeans segmentation into 5 user profiles, and executive-ready dashboards - cutting manual analysis time by **~60%** and directly informing commercial tier decisions. Separately, I ran systematic **ablation studies** across transformer block unfreezing depth (4, 6, 8 layers), learning rate schedules, and augmentation strategies, proving that disciplined context-layer control beats brute-force re-training for **out-of-distribution generalization**.

Arango's bet that **graph-based Contextual AI** is the missing layer between enterprise data and LLM reliability is the exact gap my work keeps hitting - and the place I want to build next. I would welcome the chance to discuss how I can contribute to that foundation this summer.

Sincerely,

AYMANE AIT DADS