

AYMANE AIT DADS

Research Intern - AI Agents & Efficiency | PyTorch · Transformer Fine-tuning · MLOps

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 10, 2026

Hiring Team - Research Intern, AI Agents & Efficiency

Microsoft · Redmond, WA, USA

Dear M365 Research Hiring Team,

Your call for research at the intersection of **AI efficiency** and **agentic systems** maps directly to work I have already shipped under real constraints: ranked **#1 on the EURECOM private leaderboard** by fine-tuning a **428M-parameter CLIP ViT-L/14** on 250K samples, where every design decision - FP16 mixed precision, GradScaler, and a 10-view test-time augmentation strategy - was motivated by the need to maximize generalization within a fixed GPU budget.

Two results speak directly to your research priorities. First, my **ablation study** across transformer block unfreezing depth (4, 6, and 8 layers) and learning rate schedules produced a counter-intuitive finding: **1-epoch fine-tuning** preserved out-of-distribution generalization better than extended re-training - the kind of efficiency insight that matters when scaling agentic inference. Second, at Orange Maroc I engineered an **end-to-end ML pipeline** over **100K+ telemetry records**, cutting manual analysis time by **~60%** and delivering segmentation outputs that drove go-to-market decisions - evidence that I build for impact, not just benchmark scores.

M365 Research sits in a rare position: publishing at top venues while shipping to hundreds of millions of users through products like Copilot, where agent efficiency is not a research toy problem but a latency and cost constraint at global scale. I would welcome the chance to discuss how my work on inference efficiency and rigorous experimentation can contribute to that pipeline - please feel free to reach me at [Aymane.Ait-dads@eurecom.fr](mailto:Ayman.Ait-dads@eurecom.fr).

Sincerely,

AYMANE AIT DADS