

AYMANE AIT DADS

Developer Technology Intern, High-Performance Databases | GPU Optimization · ML Pipelines · PyTorch

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 01, 2026

Hiring Team - Developer Technology Intern, High-Performance Databases - Fall 2026

NVIDIA Corporation · Santa Clara, USA

Dear NVIDIA DevTech Hiring Team,

The DevTech team is looking for someone who can **profile, analyze bottlenecks, and squeeze every single bit of performance** out of GPU-accelerated workloads - that precise discipline is what I applied when fine-tuning **CLIP ViT-L/14 (428M parameters)** on **250K samples** under T4 GPU memory constraints, combining FP16 mixed precision with GradScaler to maximize throughput without sacrificing model quality, ultimately ranking **#1 on the EURECOM private leaderboard** across 25+ generator types.

Two additional proof points map directly to this role. First, to address the challenge of **data ingestion pipelines constraining AI workload performance**, I engineered an end-to-end **ML pipeline processing 100K+ telemetry records** at Orange Maroc - from raw ingestion to business-ready output - cutting manual analysis time by **~60%**. Second, my ablation experiments on transformer block **unfreezing depth (4/6/8 layers)** and learning rate schedules demonstrate the kind of architecture-level profiling and optimization thinking that DevTech performance engineers apply daily.

NVIDIA's cuDF and RAPIDS stack are pushing GPU-accelerated data analytics into territory where CPU-based tools simply cannot compete - and that is exactly the performance frontier I want to work on. I would welcome the chance to discuss how my background fits the Fall 2026 DevTech internship.

Sincerely,

AYMANE AIT DADS