

AYMANE AIT DADS

Generative AI Intern | LLMs · RAG Pipelines · Transformer Fine-Tuning

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 11, 2026

Hiring Team - Generative AI Intern

Openkyber · Remote / International

Dear Openkyber Hiring Team,

Your requirement to **fine-tune and optimize models using techniques like LoRA** and prompt engineering maps directly to what I have already done in production: I fine-tuned **CLIP ViT-L/14 (428M parameters)** on 250K images spanning 25+ generative model types - running systematic ablation studies across transformer block unfreezing depth, learning rate schedules, and augmentation strategies - and ranked **#1 on the EURECOM private leaderboard**. That is not a notebook experiment; it is disciplined model optimization under real GPU constraints.

Two more JD requirements I can hit from day one: first, **building and optimizing RAG pipelines** - my NLP work includes a full transformer fine-tuning pipeline from TF-IDF baseline through word2vec embeddings to BERT-style fine-tuning, gaining **12 percentage points** over baseline, the exact stack that feeds document intelligence and retrieval systems. Second, **deploying AI via REST APIs and microservices** - at Orange Maroc I owned an end-to-end ML pipeline over **100K+ telemetry records**, delivering outputs directly into executive dashboards used for go-to-market decisions, not handed off to an engineer to productionize.

Openkyber's focus on enterprise-grade LLM integration - connecting models like **GPT, Claude, and LLaMA** to real business systems through vector databases and microservices - is precisely the kind of production-grade AI work I am looking to contribute to. I would welcome the chance to walk through my work in a short call.

Sincerely,

AYMANE AIT DADS