

AYMANE AIT DADS

Agentic AI Engineering Intern - Cybersecurity | PyTorch · Transformer Fine-tuning · MLOps
Aymane.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 01, 2026

Hiring Team - Agentic AI Engineering Intern - Cybersecurity

Qualys · Foster City, CA, USA

Dear Qualys Hiring Team,

Your requirement to **identify and resolve conflicts or ambiguities across tool definitions that lead to incorrect agent behavior** maps directly to work I know well: building systems that must make consistent, reliable decisions under ambiguity. Fine-tuning **CLIP ViT-L/14 (428M parameters)** across **25+ generator types** to rank **#1 on the EURECOM leaderboard** required exactly that - designing inference logic that selects the right classification path consistently across out-of-distribution inputs, with ablation experiments to prove where and why the system failed before it succeeded.

On the **evaluation harness** side: I designed **10-view test-time augmentation** and ran structured ablation experiments across transformer block unfreezing depth (4, 6, and 8 layers) and learning rate schedules to track model behavior changes systematically - directly analogous to building frameworks that measure **tool selection accuracy** over time. On the **agent orchestration and pipeline** side: at Orange Maroc I owned an end-to-end **ML pipeline** over **100K+ telemetry records**, reducing manual analysis time by **~60%** and delivering outputs used in real executive decisions - evidence that I build for production, not just experimentation.

Qualys operates one of the largest cloud-native security platforms in the world, and the shift toward autonomous agents acting on security APIs is precisely where rigorous capability boundary design matters most - a system that escalates incorrectly is as dangerous as one that acts without authorization. I would welcome the chance to discuss how I can contribute to this directly.

Sincerely,

AYMANE AIT DADS