

AYMANE AIT DADS

Edge AI Engineering Intern | PyTorch · Model Optimization · Edge Deployment

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 13, 2026

Hiring Team - Intern Software Engineer - Edge AI Engineering

Roofline · Europe

Dear Roofline Hiring Team,

Roofline's mission to **simplify AI deployment at the edge** by transforming how models run on **diverse hardware** maps directly to work I have already done under real constraints: fine-tuning a **428M-parameter CLIP ViT-L/14** on 250K images with FP16 mixed precision and GradScaler to fit within a T4 GPU budget - then designing inference pipelines robust enough to rank **#1 on the EURECOM leaderboard** against 25+ generator types.

On the optimization side, I ran systematic ablation studies over transformer block **unfreezing depth (4/6/8 layers)** and learning rate schedules, proving that selective layer unfreezing preserves generalization - precisely the kind of efficiency-accuracy tradeoff that matters when deploying on constrained edge hardware. Separately, I built an **end-to-end ML pipeline** processing **100K+ telemetry records** at Orange Maroc using **scikit-learn and Python**, reducing manual analysis time by **~60%** - showing I can ship production-grade systems, not just research prototypes.

Roofline sits at the intersection of model optimization and hardware-aware deployment - a space where most teams choose one or the other. That combination is exactly where I want to build expertise. I would welcome the chance to discuss how I can contribute to your edge AI engineering work this summer.

Sincerely,

AYMANE AIT DADS