

AYMANE AIT DADS

AI Research Intern - Agentic AI | PyTorch · Transformer Fine-tuning · LLM Optimization

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

April 30, 2026

Hiring Team - AI Research Intern - Agentic AI

Samsung Electronics · Staines-upon-Thames, United Kingdom

Dear Samsung Research UK Hiring Team,

Your posting calls for someone who can **finetune and deploy LLMs** using optimization techniques including **supervised finetuning** - and ship that work into real users' hands. I ranked **#1 on the EURECOM leaderboard** by fine-tuning a **428M-parameter CLIP ViT-L/14** on 250K samples across 25+ generator types, applying FP16 mixed precision and GradScaler to make a large transformer viable under real GPU constraints - not just a research artifact, but an optimized, deployable system.

On the research rigor your role demands: I ran **systematic ablation studies** over transformer block unfreezing depth (**4, 6, and 8 layers**), learning rate schedules, and epoch counts - proving that 1-epoch fine-tuning generalizes better than extended re-training, a finding directly relevant to **on-device model efficiency**. Separately, at Orange Maroc I owned an **end-to-end ML pipeline** across **100K+ telemetry records**, delivering segmentation outputs that directly informed commercial decisions - the kind of research-to-product translation SRUK describes.

Samsung's push to bring **agentic AI and LLMs onto mobile platforms** - running autonomously in users' pockets rather than in a data center - is exactly the constraint-driven research problem I want to work on. I would welcome the chance to discuss how my optimization and fine-tuning experience maps to your on-device agent work.

Sincerely,

AYMANE AIT DADS