

AYMANE AIT DADS

Generative AI Engineer Intern | PyTorch · Transformer Fine-tuning · LLM Deployment

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

April 30, 2026

Hiring Team - SAP Intern - Generative AI Engineer

SAP · Singapore

Dear SAP GenAI CoE Hiring Team,

Your listing calls for hands-on work with **Large Language Model deployments** - experimenting with frameworks like **vLLM** and **LiteLLM** to build scalable, production-ready Generative AI services. That maps directly to what I do: I fine-tuned a **428M-parameter CLIP ViT-L/14** transformer on 250K samples, ran systematic ablation studies across unfreezing depth and learning rate schedules, and ranked **#1 on the EURECOM private leaderboard** - not by brute-forcing compute, but by understanding how large models behave under real constraints.

Two proof points directly relevant to your stack: First, building **scalable, reusable services** - at Orange Maroc I engineered an **end-to-end ML pipeline** over **100K+ telemetry records**, cutting manual analysis time by **~60%** and delivering segmentation outputs that directly informed go-to-market decisions. Second, optimizing for **production system performance** - I applied **FP16 mixed precision with GradScaler** on a T4 GPU budget and designed **10-view test-time augmentation** to maintain inference robustness under distribution shift, the kind of efficiency tradeoff that matters when maintaining high-availability platforms.

SAP's GenAI CoE sits at the intersection of enterprise scale and cutting-edge AI infrastructure - co-developing an **LLM Platform-as-a-Service** used across SAP's global operations is exactly the production-grade scope I want to work in, not a sandbox. I would welcome the chance to discuss how I can contribute to your Singapore team starting June 2026.

Sincerely,

AYMANE AIT DADS