

AYMANE AIT DADS

ML Platform Fullstack Intern | PyTorch · MLOps · Transformer Fine-tuning

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 05, 2026

Hiring Team - Internship, Fullstack Software Engineer, Machine Learning Platform, AI Infrastructure (Fall 2026)

Tesla · Palo Alto, United States

Dear Tesla ML Platform Hiring Team,

Your team is building tools to **schedule, manage, and monitor machine learning experiments, data pipelines, and artifacts** at large-scale GPU clusters - exactly the infrastructure layer I have been working closest to. Fine-tuning **CLIP ViT-L/14 (428M parameters)** on 250K samples under real GPU constraints forced me to own every layer from data ingestion to inference optimization, earning the **#1 rank on the EURECOM private leaderboard** across all submissions.

To keep training tractable on a T4 GPU, I implemented **FP16 mixed precision with GradScaler** and ran systematic ablation experiments across transformer block unfreezing depth (4, 6, and 8 layers), learning rate schedules, and augmentation strategies - the kind of experiment lifecycle tracking your platform is designed to support. At **Orange Maroc**, I independently owned an **end-to-end ML pipeline over 100K+ telemetry records**, from raw ingestion to **Power BI dashboards** presented to senior management, cutting manual analysis time by roughly **60%**.

Tesla's Autopilot and FSD teams retrain neural networks continuously at a scale where platform reliability is not a support function - it is the bottleneck. I want to contribute to the tooling that removes that bottleneck. I would welcome the chance to discuss how my background maps to the team's current priorities.

Sincerely,

AYMANE AIT DADS