

AYMANE AIT DADS

AI Compiler Intern | PyTorch · FP16 · Transformer Inference Optimization

Ayman.Ait-dads@eurecom.fr | +33 7 60 92 50 93 | linkedin.com/in/aymane-ait-dads | Sophia Antipolis, France

May 09, 2026

Hiring Team - Internship, Software Engineer, AI Compiler (Summer 2026)

Tesla · Palo Alto, United States

Dear Tesla AI Compiler Hiring Team,

Your team is looking for someone who can help **extract maximum performance out of hardware** for neural network inference - the exact constraint I engineered around when fine-tuning **CLIP ViT-L/14 (428M parameters)** on a single T4 GPU, applying **FP16 mixed precision with GradScaler** to stay within budget while ranking #1 on the EURECOM leaderboard across 250K training samples.

Deploying **production models** at scale requires understanding the gap between training and real-world inference. I designed **10-view test-time augmentation** to harden out-of-distribution robustness, and ran systematic ablation studies across **transformer block unfreezing depth (4, 6, and 8 layers)** to prove that 1-epoch fine-tuning generalizes better than extended re-training - the kind of inference-time analysis your compiler/runtime co-design depends on. Separately, I shipped an **end-to-end ML pipeline over 100K+ telemetry records** at Orange Maroc, reducing manual analysis time by 60%, demonstrating I can build systems that hold up under real operational constraints.

Tesla's inference stack runs inside millions of vehicles and Optimus - there is no comparable production scale in AI deployment, and the compiler layer is where hardware and model architecture decisions converge in ways that actually matter. I would welcome the chance to contribute to that stack on-site this summer; please find my CV attached.

Sincerely,

AYMANE AIT DADS